

# Himanshu Choudhary

✉ himanshuchoudhary448@gmail.com ☎ (917)378-7543 🔗 in/himanshu-choudhary-b39b62249 🌐 github.com/himanshuchoudhary448

## SUMMARY

Software Engineer with over 5 years of experience in application development, cloud infrastructure, and DevOps across AWS and Azure. Skilled in Java, Python, TypeScript, CI/CD, Docker, Kubernetes, and Terraform, delivering scalable, cost-efficient enterprise solutions.

## EXPERIENCE

### Senior Software Engineer

People Inc.

September 2024 - Present, New York, US

- Designed a unified AWS Cost Monitoring Dashboard using CUR, S3, Athena, Glue, and QuickSight for real-time cost insights across 90 accounts.
- Developed an automated asset recovery and monitoring system with Lambda, DynamoDB, CloudWatch, and Oomnitza API, enabling real-time SLA and latency alerts.
- Developed modular Terraform pipelines executed via Jenkins and Docker to provision infrastructure on AWS, enabling secure, gated deployments across 5 regions with consistent, compliant configurations.

### Software Engineer

Morgan Stanley

November 2023 - September 2024, New York, US

- Managed VMware GemFire and Azure Redis upgrades, ensuring seamless cross-platform compatibility for Java, .NET, and Python applications.
- Implemented CI/CD pipelines with Bitbucket, Jenkins, and GitHub Actions after migrating 10 repositories from Azure DevOps, boosting delivery speed and consistency.
- Deployed high-performance Redis clusters on OpenShift with Kubernetes and Helm, integrated Argo CD for GitOps-driven deployments, enabling scalable caching and reducing response times by 20 ms.

### Software Engineer

Jefferies

November 2022 - September 2023, New Jersey, US

- Engineered a responsive, component-based financial dashboard using React.js, TypeScript, and Tailwind CSS, leveraging RESTful APIs for real-time data integration. Automated data extraction workflows, eliminating 8 hours of manual reporting weekly for analysts.
- Deployed 12 containerized applications on Amazon ECS, implementing monitoring with AWS CloudWatch and Grafana to optimize performance and reliability.
- Automated infrastructure provisioning with Python, AWS CloudFormation, and SDKs, cutting deployment timelines by 30% through consistent, repeatable setups.

### Software Engineer

AT&T

August 2021 - October 2022, Texas, US

- Architected a serverless data pipeline in Python on AWS using Glue, S3, and Athena for near real-time analysis of 10GB+ daily data.
- Built CI/CD pipelines in Groovy with Jenkins to package Java apps via Maven, deploy to Databricks, run parameterized jobs, and store computed Parquet outputs in Azure Blob Storage, reducing manual effort by 80%.
- Led analytics for 5+ projects on Databricks, optimizing Python and Spark workflows and cluster performance for faster, cost-efficient processing.

### Software Engineer

Chubb

August 2020 - July 2021, New Jersey, US

- Implemented Azure Logic Apps and Azure Functions with enhanced monitoring via Azure Monitor, improving application responsiveness and reliability, reducing incident resolution time by 25%.
- Engineered a microservices-based Java policy management system on Azure App Service, containerized with Docker and deployed on AKS. Implemented CI/CD with Azure DevOps, delivering 30% faster deployments and scalability.
- Built a Python-based premium optimization model with Azure Functions to adjust rates by risk factors, increasing underwriting profitability by 15%.

## EDUCATION

### Master of Science, Computer Science

New York Institute of Technology • New York, NY • 2020

### Bachelor of Engineering, Computer Science

Rajiv Gandhi Proudlyogiki Vishwavidyalaya • Madhya Pradesh, India • 2018

## CERTIFICATIONS

### AWS Certified Solutions Architect – Associate

Amazon Web Services (AWS) • 2025

## SKILLS

**Front-end Technologies:** HTML5, CSS3, Typescript, JavaScript Back-end Technologies: Java, Python, Spark, Scala, C#, Spring, Bash, Groovy

**Databases:** VMware GemFire, Azure Redis Cache, MongoDB, Dynamo DB, Delta Lake, MySQL, Oracle (PL/SQL)

**Big Data Technologies:** PySpark, Hadoop, Hive, Scala, HDFS, Pig, SQLite, Databricks

**DevOps & Cloud:** AWS, Azure, Jenkins, Terraform, Azure DevOps, GitHub Actions, Bitbucket, JFrog, Docker, Kubernetes, Ansible